

Facial Recognition using Eigenface

ME 766 - High Performance Scientific Computing

Shubham Anand
160050060

Sudhanshu Raj
160050041

Ankit
160050044

Simranpreet Singh
160050111

1. INTRODUCTION

- The task of face recognition is to recognise a person from a given image. This is done by matching some set of feature derived from the given image with those existing in our database.
- There are some patterns that in all the input images. These patterns are being analyzed to identify one specific image from a set of input images.
- We compute characteristic features known as eigenfaces and calculate the eigen-coefficients for each image using these eigenfaces.
- Eigenface based facial recognition is used to identify a particular face from a large database. Our aim is to analyse the algorithm and parallelise various segment of the program to enhance the speed-up (like matrix related operations) by using viennaCL library.

2. SAMPLE INPUT/OUTPUT FORMAT

- The database consists of images of 9 people and corresponding to each person there are 30 images. In total there are 270 images.
- These are gray scale images of size 168 x 192.
- The output of the code is the index of the image to which the input image is matched most closely to the above 270 images.
- The number of images which is 270 is not fixed. For analysing the timing of the code, we these number is varied from 180 to 540.

3. ALGORITHM

- For each gallery image, we compute the eigen-coefficients. We then store the eigen-coefficients and the identity of the person in a database.
- During the testing phase, we are given a probe image. We then compute the eigen coefficients for this image.
- We compare the above computed eigen-coefficients with eigen-coefficients stored in our database and find the closest match in terms of sum squared error.

4. LANGUAGE AND LIBRARY USED

- C++ for the base code
- CUDA for parallelization
- OpenMP for parallelization
- ViennaCL Library for computing eigenvalues/vectors and matrix multiplication

5. LIBRARY FUNCTIONS USED

- **prod** - For multiplying two matrices
- **copy** - For copying a data structure from host to device and vice-versa
- **qr_method_sym** - Calculating the eigenvalues and eigenvectors of a symmetric matrix
- **trans** - For transposing a matrix

6. TIMING ANALYSIS (for Input consisting of 270 images)

We used :

- ❑ **nvprof** for profiling
- ❑ **time** (command) for calculating the overall time for execution
- ❑ **nvvp** for visualising the performance and profiling

- Overall Time for Serial Execution : **82.512 s**
- Overall Time for Parallel Execution (OpenMP) using 8 processors : **26.02 s**
- Overall Time for Parallel Execution (CUDA) : **2.135 s**
- Speedup for CUDA Execution = $82.512/2.135 = 40.052$
- Speedup for OpenMP Execution = $82.512/26.02 = 3.286$

Comparison of self-defined vs library Matrix Multiplication Function-

- Time taken using self-defined function = 1.72 s
- Time taken using library function = 0.128 s

7. PROFILING

The profiling as done by nvprof is presented below --

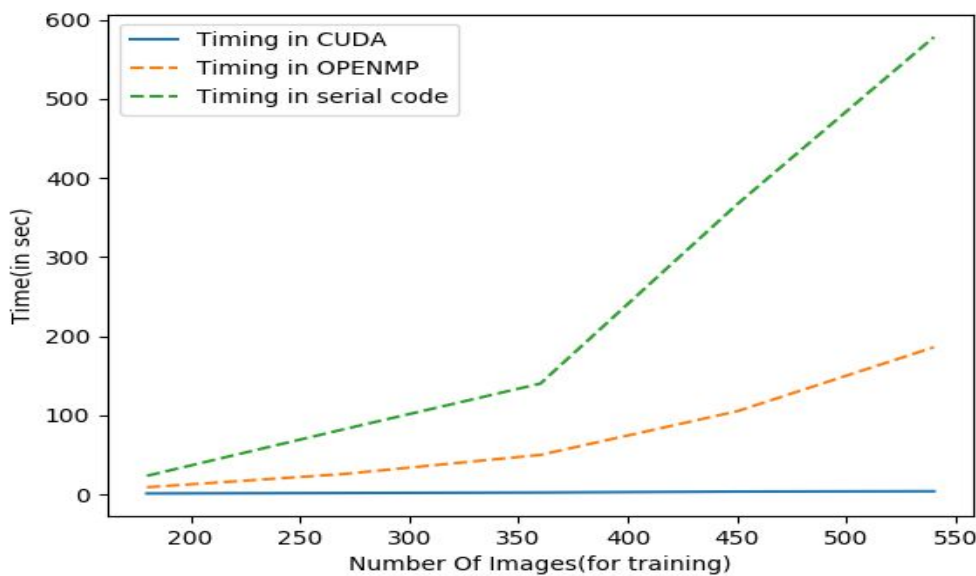
```
$ nvprof ./a.out

==31566== NVPROF is profiling process 31566, command: ./a.out
==31566== Profiling application: ./a.out
==31566== Profiling result:
   Type Time(%) Time      Calls      Avg      Min      Max Name
GPU activities: 65.14% 178.08ms    4 44.521ms 5.5187ms 65.568ms void viennacl::linalg::cuda::matrix_matrix_row_row_row_prod_
   15.74% 43.029ms    248 173.50us 608ns 13.373ms [CUDA memcpy HtoD]
   9.36% 25.592ms    186 137.59us 640ns 15.064ms [CUDA memcpy DtoH]
   3.07% 8.3939ms     1 8.3939ms 8.3939ms 8.3939ms cuda_mean(double*, double*)
   1.95% 5.3220ms    178 29.899us 2.3680us 70.913us void viennacl::linalg::cuda::house_update_A_left_row_major_kernel
   1.80% 4.9169ms     1 4.9169ms 4.9169ms 4.9169ms void viennacl::linalg::cuda::trans_kernel
   1.41% 3.8666ms     11 351.51us 2.2720us 892.46us void viennacl::linalg::cuda::matrix_row_assign_kernel
   0.71% 1.9295ms    178 10.839us 10.240us 13.056us void viennacl::linalg::cuda::house_update_QL_row_major_kernel
   0.69% 1.8741ms    178 10.528us 9.6960us 12.384us void viennacl::linalg::cuda::house_update_A_right_row_major_kernel
   0.08% 207.27us    178 1.1640us 960ns 1.6640us void viennacl::linalg::cuda::copy_col_row_major_kernel
   0.06% 173.22us    32 5.4130us 2.6560us 7.8720us void viennacl::linalg::cuda::givens_next_row_major_kernel
   0.00% 8.4160us     6 1.4020us 1.0560us 1.7600us void viennacl::linalg::cuda::vector_assign_kernel
   0.00% 1.7600us     1 1.7600us 1.7600us 1.7600us void viennacl::linalg::cuda::bidiag_pack_row_major_kernel
   0.00% 1.6320us     1 1.6320us 1.6320us 1.6320us void viennacl::linalg::cuda::matrix_row_diagonal_assign_kernel
API calls: 58.11% 213.23ms    434 491.31us 6.0650us 69.149ms cudaMemcpy
   38.04% 139.59ms    18 7.7552ms 6.5540us 135.61ms cudaMalloc
   1.78% 6.5198ms    16 407.49us 5.6200us 1.9130ms cudaFree
   1.67% 6.1199ms    769 7.9580us 6.3740us 35.356us cudaLaunch
   0.19% 702.67us   4793 146ns 117ns 771ns cudaSetupArgument
   0.12% 422.22us    94 4.4910us 543ns 167.33us cuDeviceGetAttribute
   0.05% 185.47us    769 241ns 157ns 1.0700us cudaConfigureCall
   0.02% 86.098us     1 86.098us 86.098us 86.098us cuDeviceTotalMem
   0.01% 52.917us     1 52.917us 52.917us 52.917us cuDeviceGetName
   0.00% 5.7180us     3 1.9060us 652ns 4.2410us cuDeviceGetCount
   0.00% 5.6800us    20 284ns 211ns 407ns cudaGetLastError
   0.00% 2.2640us     2 1.1320us 624ns 1.6400us cuDeviceGet
```

Profiling Visualisation using `nvvp --`



Timing Comparison of **Serial**, **OpenMP** and **CUDA** execution --



8. CONCLUSION

We observe that the speedup achieved through CUDA execution (**40.052**) is far more better than the speedup achieved through openMP execution (**3.286**). This is due to the fact that there are more number of compute elements present in GPU than what can be achieved through the processors.

Using library functions (ViennaCL) also has a significant impact on the speedup. We have demonstrated this using timing analysis and comparing the two instances ie. with our own Matrix Multiplication function and one provided by the library.